

2. The psql Client

Introduction to PostgreSQL



PostgreSQL

AGENDA

- Introduction to the psql client
- Configuring client access pg_hba.conf
- Command line parameters
- Meta-commands
- Input and Output



PSQL

- The *psql* client is a command-line interface (CLI) used to interact with the PostgreSQL database server.
- It is the primary tool for administration, querying, and managing PostgreSQL databases.
 - *Interactive Query*: query, manipulate, and manage database data by running SQL commands.
 - *Administration*: perform administrative tasks like managing users, and setting permissions.
 - *Script Execution*: can run SQL scripts to automate batch processing of database operations.



CONNECTION OPTIONS

- -h, --host
 - Specifies the host name or IP address of the PostgreSQL server to connect to. Defaults to a local connection if not specified.
 - *psql -h localhost or psql --host=192.168.1.100*
- -p, --port
 - Specifies the TCP port number the server is listening on. The default port is 5432.
 - *psql -p 5432 or psql --port=5433*
- -U, --username
 - Specifies the PostgreSQL username to connect as.
 - *psql -U postgres or psql --username=myuser*

CONNECTION OPTIONS

- -d, --dbname
 - Specifies the name of the database to connect to.
 - If omitted, psql connects to the database with the same name as the user.
 - *psql -d mydatabase or psql --dbname=mydatabase*
- -W, --password
 - Forces psql to prompt for a password before connecting. This ensures a password is always requested.

OTHER CONNECTION OPTIONS

- Often, we want to not open an interactive shell but just perform a specific operation.
- `--l, --list`
 - Lists all databases in the server
 - `psql -l`
- `-c, --command`
 - Executes the specified SQL command and then exits.
 - `psql -c "SELECT version();"`
- `-f, --file`
 - Executes commands from the specified file and then exits. This is useful for running SQL scripts.
 - `psql -f myscript.sql` or `psql --file=myscript.sql`

OTHER CONNECTION OPTIONS

- -F, --field-separator
 - Specifies the field separator for unaligned output format, overriding the default tab character.
 - *psql -F ',' to use a comma as the separator.*
- -A, --no-align
 - Switches to unaligned output mode, which prints columns without any padding, making it suitable for script processing.
- -t, --tuples-only
 - Outputs only the data rows, without headers or other formatting information. Useful for scripting or exporting data.
 - *psql -t -c "SELECT * FROM mytable;"*

OTHER CONNECTION OPTIONS

- -q, --quiet
 - Runs psql in quiet mode, suppressing some output such as welcome messages and command status.
- --echo-queries
 - Echoes SQL queries to standard output before executing them. Used for debugging.
- --no-password
 - Prevents psql from prompting for a password. It fails if a password is required

OTHER CONNECTION OPTIONS

- `--single-transaction`
- Executes all commands in a single transaction. If an error occurs, all changes are rolled back.
- *`psql --single-transaction -f myscript.sql`*
- `-L, --log-file`
- Logs all input to the specified file, useful for debugging or keeping a record of commands executed during a session.
- *`psql -L mylogfile.log`*

LAB 2-1

- The lab description and documentation is in the Lab directory in the class repository



PSQL META-COMMANDS

- Meta-commands begin with a backslash (\) and provide a ways to interact with PostgreSQL other than by SQL commands.
 - These commands are processed by `psql` itself, not the database server.
 - Many commands do not require any interaction with the serve, like `\q` and `\c database`
 - Those require information from the server, like `\dt`, are parsed into the appropriate SQL command and sent to the server.
 - The generated `\dt` query fetches information about the tables, from the `pg_class` and `pg_namespace` system catalogs.
 - `psql` formats the results of the query into human-readable tabular format which is displayed in `psql` console output

PSQL META-COMMANDS

- The \? command lists all the meta-commands available
 - Sample commands are listed for each group
- The main command groups we will be looking at are:
 - Database Listing Commands
 - Database Object Description Commands
 - Session Management Commands
 - Query Execution Control Commands
 - File and Script Execution Commands
 - Environment Variable Management Commands
 - Administrative Commands

```
student=> \?
General
\copyright          show PostgreSQL usage and distribution terms
\crosstabview [COLUMNS] execute query and display result in crosstab
\errverbose         show most recent error message at maximum verbosity
\g [(OPTIONS)] [FILE] execute query (and send result to file or |pipe);
                   \g with no arguments is equivalent to a semicolon
\gdesc             describe result of query, without executing it
\gexec             execute query, then execute each value in its result
\gset [PREFIX]     execute query and store result in psql variables
\gx [(OPTIONS)] [FILE] as \g, but forces expanded output mode
\q                 quit psql
\watch [SEC]       execute query every SEC seconds

Help
\? [commands]      show help on backslash commands
\? options          show help on psql command-line options
\? variables        show help on special variables
\h [NAME]           help on syntax of SQL commands, * for all commands
```

DATABASE LISTING COMMANDS

- Displays information about the database and its objects, such as tables, indexes, schemas, and roles.
 - \l or \list: Lists all databases on the server.
 - \dt: Lists all tables in the current database.
 - \dv: Lists all views.
 - \di: Lists indexes.
 - \df: Lists functions.
 - \dn: Lists schemas.
 - \du: Lists roles/users

DATABASE OBJECT DESCRIPTION COMMANDS

- Provides detailed descriptions of specific database objects.
 - \d [object_name]: Describes a table, view, index, or sequence.
 - \d+ [object_name]: Provides a more detailed description, including additional metadata.
 - \df+ [function_name]: Shows detailed information about functions, including argument types and return type.

SESSION MANAGEMENT COMMANDS

- Manages database connections and the session environment, including connecting to different databases or switching users.
 - \c [database_name] or \connect [database_name]: Connects to a different database.
 - \conninfo: Displays information about the current connection.
 - \password [user]: Changes the password of a database user.

QUERY EXECUTION CONTROL COMMANDS

- Controls how SQL commands and queries are executed and displayed.
 - \timing: Toggles timing of commands, useful for performance analysis.
 - \watch [seconds]: Re-executes the last command every specified number of seconds.
 - \g [filename]: Sends query results to a file or program.

FILE AND SCRIPT EXECUTION COMMANDS

- Allows users to execute SQL commands from files or save query results to files.
 - \i [file_name]: Executes commands from a specified SQL file.
 - \o [file_name]: Redirects query output to a file.
 - \ir [file_name]: Executes SQL file, but reads it relative to the current working directory.
 - \copy: Used to copy data from a file to a table

FILE AND SCRIPT EXECUTION COMMANDS

- Allows users to execute SQL commands from files or save query results to files.
 - \i [file_name]: Executes commands from a specified SQL file.
 - \o [file_name]: Redirects query output to a file.
 - \ir [file_name]: Executes SQL file, but reads it relative to the current working directory.
 - \copy: Used to copy data from a file to a table

ENV VARIABLE MANAGEMENT COMMANDS

- Manages variables and the environment settings in the psql session.
 - `\set [variable] [value]`: Sets a psql variable.
 - `\unset [variable]`: Unsets a psql variable.
 - `\echo :[variable]`: Displays the psql variable contents
 - `\prompt [text] [variable]`: Prompts the user for input and stores it in a variable.

ADMINISTRATIVE COMMANDS

- Used for administrative tasks
 - \dconfig: Shows configuration parameters.
 - \dconfig+: Shows configuration parameters with comments and source.
 - \! [cmd]: Executes a shell command
 - \h [SQL command]: Provides help for SQL commands.

LAB 2-2

- The lab description and documentation is in the Lab directory in the class repository



SECURITY CONFIGURATION

- pg_hba.conf
 - PostgreSQL Host-Based Authentication
 - A file is used by PostgreSQL to control client authentication.
 - It specifies which users can connect to which databases from which hosts, and what authentication methods they use.
 - The location of the file being used by the running instance can displayed using:
 - *SHOW hba_file;*



SECURITY CONFIGURATION

- The file is a set of records, one per line.
- Each record specifies:
 - A connection type
 - A client IP address range if needed
 - A database name
 - A username,
 - The authentication method to be used for connections matching these parameters.




SECURITY CONFIGURATION

- The first record that matches all the connection parameters is used to perform authentication.
 - If authentication fails for one record, other records are not considered.
 - If no record matches, access is denied.




CONNECTION TYPES

- local
 - Connections using UNIX domain sockets.
 - Used for connections from the same machine where PostgreSQL is running.
 - Commonly used with authentication methods like peer that map PostgreSQL users with system users
- host
 - Connections attempts made using TCP/IP for either SSL or non-SSL connections
- hostssl
 - Connections that are secured using SSL
 - Ensures encrypted communication between the client and server.



CONNECTION TYPES

- hostnossl
 - Matches connection attempts made over TCP/IP that do not use SSL.
- hostgssenc
 - Connections made using TCP/IP with GSSAPI encryption
 - Used when Kerberos authentication is configured for secure connections
- hostnogssenc
 - Matches connection attempts made over TCP/IP that do not use GSSAPI.



DATABASE SPECIFICATION

- Specifies which database name(s) a record matches.
 - *all* matches all databases
 - *sameuser* matches if the database has the same name as the user
 - *replication* specifies that the record matches if a physical replication connection is requested
 - If the database name starts with a slash (/), the remainder of the name is treated as a regular expression
- Multiple database names and/or regular expressions can be supplied by separating them with commas




USER SPECIFICATION

- Specifies which database user name(s) this record matches with
 - *all* for all users
 - Specific PostgreSQL usernames
 - A regular expression starting with a /
 - OS level groups specified by +groupname
 - A comma delimited list of names and regular expressions



CONNECTION ADDRESSES

- Not required for non-network connections
 - Like local Unix sockets
- IP addresses that connections are allowed from
 - Can also be written as an IP address and netmask



AUTHENTICATION TYPES

- Some common authentication methods
 - *trust*: allows anyone that can connect to the server to login as any user they wish, without the need for authentication..
 - *reject*: reject the connection unconditionally - useful for “filtering out” certain hosts.
 - *scram-sha-256*: SCRAM-SHA-256 authentication to verify the password.
 - *md5*: SCRAM-SHA-256 or MD5 authentication to verify the password.
 - *peer*: uses client's operating system user name from the operating system, only available on local connection




SECURITY CONFIGURATION

- postgresql.conf file
 - Specifies the connections that are allowed
- listen_addresses parameter
 - Specifies the TCP/IP address(es) on which the server is to listen for connections from client applications.
 - The special entry * corresponds to all available IP interfaces.
 - The entry 0.0.0.0 allows all IPv4 addresses and :: allows all IPv6 addresses.
 - If the list is empty, only Unix-domain sockets can be used to connect to it.
 - The default is *localhost*, which allows only local TCP/IP “loopback” connections to be made.




SECURITY CONFIGURATION

- `unix_socket_directories` parameter
 - Specifies the directory of the Unix-domain socket(s) on which the server is to listen for connections from client applications.
 - Multiple sockets can be specified
 - The default value is normally `/tmp`
 - On Windows, the default is empty



SECURITY CONFIGURATION

- `unix_socket_permissions` parameter
 - Sets the access permissions of the Unix-domain socket(s).
 - The default permissions are 0777, meaning anyone can connect.
 - This parameter is irrelevant on systems, that ignore socket permissions



WALK THROUGH

- Exploring the configuration files



End Module



PostgreSQL