

4. Managing Tables and Views

Introduction to PostgreSQL



PostgreSQL

AGENDA

- Basic table creation
- Table constraints
- Table relationships
- Referential integrity
- Designing tables well



TABLE MANAGEMENT

- Tables and related objects are created in PostgreSQL using standard SQL commands
- Tables are created in a schema in a database
 - If no schema is specified, then the table is created in the default public schema
- Tables are owned by the user that created them
 - Access privileges are given by the owner to other users with GRANT



TABLE CREATION

- Created using CREATE TABLE
 - List tables with \dt
 - List table details with \d <tablename>
- Delete tables with DROP TABLE
 - All data in the table is lost
 - Tables can only be dropped by the owner or user with permission to drop the table

```
CREATE TABLE weather (  
    city          varchar(80),  
    temp_lo       int,          -- low temperature  
    temp_hi       int,          -- high temperature  
    prcp          real,         -- precipitation  
    date          date  
);  
  
CREATE TABLE cities (  
    name          varchar(80),  
    location       point  
);
```

```
lab4=# DROP TABLE cities, weather;  
lab4=# \dt  
        List of relations  
 Schema | Name   | Type  | Owner  
-----+-----+-----+-----  
(0 rows)
```

CREATE TABLE AS

- Create new table by copying the data and structure of the existing table.
 - WITH NO DATA only copies the structure
 - Otherwise the data is copied
- SELECT statements can be used to populate the new table with a subset of the original data

```
lab4=# SELECT * FROM employees;
```

id	name	department	salary
1	Alice	Sales	50000
2	Bob	Marketing	55000
3	Charlie	Sales	60000
4	David	IT	65000
5	Eve	HR	48000

(5 rows)

```
lab4=# CREATE TABLE ecopy AS TABLE employees;
```

```
lab4=# SELECT * FROM ecopy;
```

id	name	department	salary
1	Alice	Sales	50000
2	Bob	Marketing	55000
3	Charlie	Sales	60000
4	David	IT	65000
5	Eve	HR	48000

(5 rows)

TABLE INHERITANCE

- In the OO view, tables represent types of object
 - We can create sub-types that inherit from base types
- The sub-type
 - Has the same columns as its parent
 - Generally has more attributes or columns
- This is used for table design rather than storing normalized data.
 - Working with data across multiple child tables can be complex
 - Often partitioning is a better solution

```
CREATE TABLE cities (  
    name      text,  
    population float8,  
    elevation  int    -- (in ft)  
);  
  
CREATE TABLE capitals (  
    state      char(2)  
) INHERITS (cities);  
  
-- Now, let's populate the tables.  
INSERT INTO cities VALUES ('San Francisco', 7.24E+5, 63);  
INSERT INTO cities VALUES ('Las Vegas', 2.583E+5, 2174);  
INSERT INTO cities VALUES ('Mariposa', 1200, 1953);  
  
INSERT INTO capitals VALUES ('Sacramento', 3.694E+5, 30, 'CA');  
INSERT INTO capitals VALUES ('Madison', 1.913E+5, 845, 'WI');  
  
SELECT * FROM cities;  
SELECT * FROM capitals;
```

TEMPLATES

- template0 and template1 are special system databases that serve as templates for creating new databases
 - template1 is used when creating new with commands like CREATE DATABASE db;
 - template0 is the pristine backup copy in case we need to reconstruct our templates

LIST OF DATABASES

Name	Owner	Encoding	Collate	Ctype	ICU Locale	Locale Provider	Access privileges
mydb	rod	UTF8	en_CA.UTF-8	en_CA.UTF-8		libc	
postgres	postgres	UTF8	en_CA.UTF-8	en_CA.UTF-8		libc	
template0	postgres	UTF8	en_CA.UTF-8	en_CA.UTF-8		libc	=c/postgres + postgres=CTc/postgres
template1	postgres	UTF8	en_CA.UTF-8	en_CA.UTF-8		libc	=c/postgres + postgres=CTc/postgres

(4 rows)

TEMPLATES

- No one can connect to template0
- However, we can connect to template1 to modify it
 - These modifications will be part of any database created from it.
- Typical modifications
 - Common extensions, such as hstore, pgcrypto, uuid-oss, PostGIS (for spatial data), or citext (for case-insensitive text)
 - *CREATE EXTENSION pgcrypto;*
 - User-defined functions or procedures that are used frequently across databases
 - *CREATE FUNCTION my_custom_function() ...*
 - Creating Common Schemas or Tables
 - *CREATE SCHEMA common_schema; CREATE TABLE common_schema.config (key text, value text);*

TEMPLATES

- Database-level configuration settings, such as `search_path`, `default_statistics_target`, or `maintenance_work_mem`
 - *ALTER DATABASE template1 SET search_path TO 'public, common_schema';*
- Creating Common Roles or Permissions
 - *CREATE ROLE readonly; GRANT SELECT ON ALL TABLES IN SCHEMA public TO readonly;*
 - *If certain indexes are always required on specific types of data or tables, these can be added*
- Adding Predefined Data
 - *INSERT INTO common_schema.config (key, value) VALUES ('default_language', 'en');*
- Any template-related configuration, such as the locale, encoding, or other database defaults
- Changing permission for the public schema to lock it down
 - *REVOKE ALL ON SCHEMA public FROM PUBLIC;*

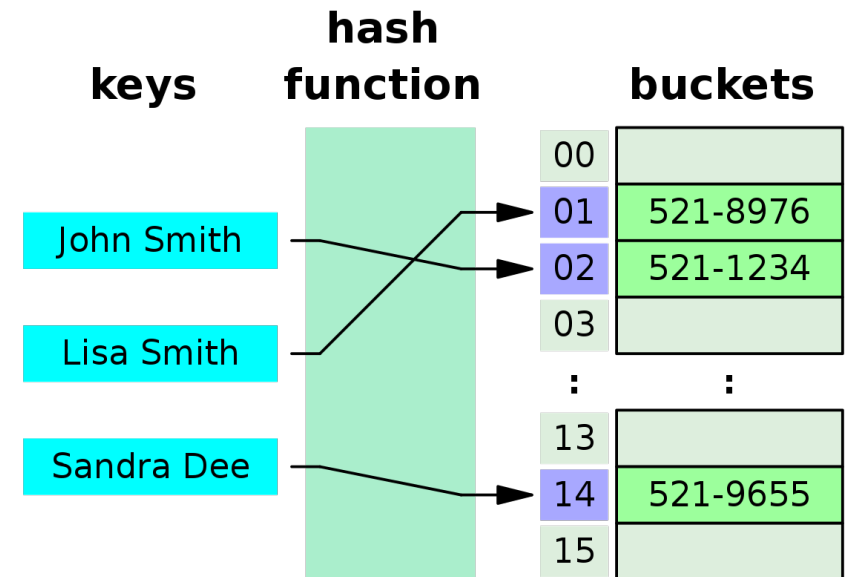
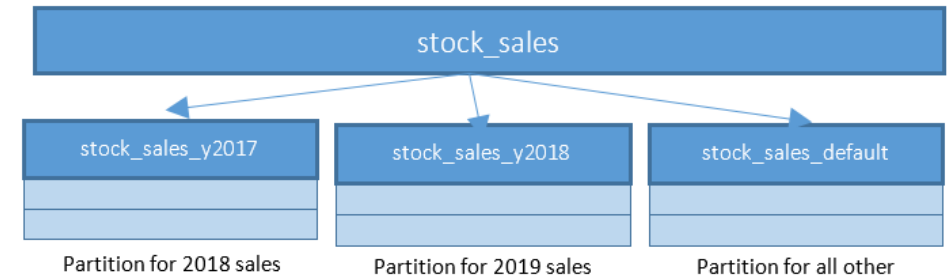
TABLE PARTITIONING

- Partitioning divides large tables into smaller, more manageable pieces
 - Different methods for partitioning are:
 - Range – data is grouped by ranges of values for a specific attribute, like transaction data
 - List – similar to range but on a categorical attribute, like county name for example
 - Hash – uses a hash c part partitioning methods. Here's an example of how to create a partitioned table using range partitioning:

```
CREATE TABLE cities (  
  name      text,  
  population float8,  
  elevation  int    -- (in ft)  
);  
  
CREATE TABLE capitals (  
  state      char(2)  
) INHERITS (cities);  
  
-- Now, let's populate the tables.  
INSERT INTO cities VALUES ('San Francisco', 7.24E+5, 63);  
INSERT INTO cities VALUES ('Las Vegas', 2.583E+5, 2174);  
INSERT INTO cities VALUES ('Mariposa', 1200, 1953);  
  
INSERT INTO capitals VALUES ('Sacramento', 3.694E+5, 30, 'CA');  
INSERT INTO capitals VALUES ('Madison', 1.913E+5, 845, 'WI');  
  
SELECT * FROM cities;  
SELECT * FROM capitals;
```

TABLE PARTITIONING

- Partitioning divides large tables into smaller, more manageable pieces
 - Range partitioning – data is grouped by ranges of values for a specific attribute, like transaction data
 - List – similar to range but on a categorical attribute, like county name for example
 - Hash – uses hashing into partitions based on hashing a specific attribute
 - Provides a more uniform distribution of data if the data is skewed along a partition attribute
 - We define the number of partitions
 - Useful with large data sets because it improves the performance of queries



TABLESPACES

- Tablespace define locations on the file system where the database objects are stored.
 - By default all data is stored in the default tablespace located in the data directory.
 - Tablespaces define other locations for storing data by mapping it to specific physical locations.
 - Created by defining a directory PostgreSQL has permission to read and write data files.
- Benefits
 - By distributing data across different physical disks, I/O performance can be improved.
 - New tablespaces on different disks can expand physical storage without moving the entire database.
 - Allows DBAs to specify where specific tables, indexes, or entire schemas are stored, making it easier to manage large and complex databases.
 - Provides a way to separate data physically.
 - Streamlines backup and recovery

TABLESPACES

- Possible issues
 - Can add complexity administration because of track of where different object are located.
 - For physical backups tablespaces must be backed up and restored to their respective locations or errors could result.
 - Logical backups don't capture tablespace information by default
- pg_default Tablespace
 - Where user defined database are stored by default unless specified otherwise.
- pg_global Tablespace
 - Used for storing shared system catalogs and objects that are shared across the entire cluster.
 - *Objects like pg_database, which keeps track of all databases in the cluster, and pg_authid, which stores role and authentication information. T*
 - Not accessible for user defined objects

VIEWS

- Virtual tables that represent the result of a query.
 - The query is stored in raw SQL format so it can be reused
 - Used to simplify complex queries, encapsulate business logic, and present data without duplicating the actual data in the underlying tables
- Limitations of Views:
 - Views execute their underlying query each time they are accessed
 - Complex views can have performance issues, especially if there are multiple joins or large datasets.
- Not all views are updateable
 - Updateable means that changes made to the view are propagated into the tables
 - Updateable must meet specific criteria, such as being based on a single table without groupings or aggregates.

VIEWS

- Updateable view requirements
 - The view must be defined by a simple SELECT statement that meets the following criteria:
 - *It selects from a single table.*
 - *It does not include any aggregation (GROUP BY, HAVING).*
 - *It does not include DISTINCT, set operations (UNION, INTERSECT, EXCEPT), or subqueries.*
 - *It does not include window functions.*
 - *It does not use LIMIT or OFFSET.*
 - *It does not contain JOIN clauses.*
 - Each column in the view must directly map to a column in the underlying base table without any expressions, computations, or column aliases that modify data types or values.
 - The underlying table should have a primary key or a unique constraint to uniquely identify rows.

MATERIALIZED VIEWS

- Similar to regular views
 - Unlike regular views, materialized views store the result set of the query on disk, allowing for faster access to the data.
 - Since the data is stored, materialized views need to be refreshed to reflect changes in the underlying tables.
 - Refreshing can be done manually or automatically.
 - Ideal for queries that are expensive to run and do not need real-time data accuracy, such as reporting or analytical queries.

DATA TYPES

- PostgreSQL supports the standard SQL data types
 - Also supports a number of complex data types
 - This will be reviewed in the documentation
- Default values can be specified
 - Either values of some computation
 - Like the 'now()' function



CREATING DATA TYPES

- PostgreSQL also allows user to create custom data types
- Common custom types
 - Composite Types: allow grouping multiple fields of different data types into a single logical unit.
 - Enum Types: static, predefined values for attributes with a limited set of possible values
 - *Excellent for enforcing referential integrity*
 - Domain Types: on existing data types but include constraints
 - *CREATE DOMAIN positive_int AS INT CHECK (VALUE > 0);*

```
CREATE TYPE full_name AS (  
    first_name TEXT,  
    last_name TEXT  
);  
  
CREATE TABLE employees (  
    employee_id SERIAL PRIMARY KEY,  
    name full_name,  
    position TEXT  
);  
  
INSERT INTO employees (name, position) VALUES  
(('John', 'Doe'), 'Software Developer'),  
(('Jane', 'Smith'), 'Project Manager');  
  
SELECT name.first_name FROM employees;
```

CUSTOM DATA TYPES

- Benefits
 - Maintain data integrity by enforcing consistent data structures and validation rules.
 - Encapsulating related fields into a single type allows for simpler queries
 - Promotes re-usability across multiple tables and applications, reducing redundancy.
 - User-defined types makes PostgreSQL highly extensible, for specific application needs.
- Considerations
 - Custom types can increase schema complexity
 - Some custom types, especially those requiring complex operations, might impact performance.
 - Might introduce compatibility challenges when interfacing with other systems or migrating data.

LAB 4-1

- The lab description and documentation is in the Lab directory in the class repository



TABLE CONSTRAINTS

- Table constraints are rules applied to columns and tables
 - Enforce data integrity, ensure consistency, and define the relationships between data.
 - Help prevent invalid data from being entered into the database
 - *Eg. Two customers with the same customer number*
 - Provide safeguards that maintain the logical correctness of the data.
 - *Eg. Invoice for a non-existent customer*
- Types of constraints
 - UNIQUE: ensures all values in a column or group of columns are unique across the table.
 - NOT NULL: Ensures that a column cannot have NULL values.
 - PRIMARY KEY: Combination of UNIQUE and NOT NULL
 - *Each table can have at most one primary key*
 - *Often a single column but can be combination of columns – called a composite key*

TABLE CONSTRAINTS

- FOREIGN KEY: Enforces a link between the data in two tables for referential integrity
 - *References the primary key or unique constraint of another table.*
 - *Prevents actions that would Invalidate the link like deleting the foreign table row if there are dependent rows in the table where the foreign key constraint is defined*
 - *Can specify actions on update or delete (e.g., CASCADE, SET NULL, RESTRICT)*
- CHECK: Enforces custom conditions that must be true for each row in the table.
 - *Similar to constraints when defining a domain.*
- EXCLUSION: Ensures that, for any two rows, the specified columns do not have overlapping values
 - *Commonly used with range and geometric data types.*
 - *Useful for scenarios like scheduling where no two time periods should overlap.*

TABLE CONSTRAINTS

- Constraints ensure that the data entered into a table meets specified criteria, maintaining data integrity and reliability.
- Constraints can also impact performance, especially on large tables with frequent updates.
 - Proper indexing and query optimization are necessary to mitigate performance degradation
- Constraints can be combined to enforce complex rules.
 - For instance, a column can be both NOT NULL and UNIQUE
- Some constraints, such as foreign keys, declared DEFERRED
 - This means they are checked at the end of a transaction rather than immediately

RELATIONAL CONSTRAINTS

- The primary use of a relational constraint
 - Defines what should happen to a row in a child table that has a foreign key (parent table)
 - Changes in the parent table can invalidate the relationship, not the data
 - *For example, the customer table has a reference to the sales person table. What happens when a sales person is fired?*
- The changes that can affect referential integrity are DELETE and UPDATE for the parent table
- CASCADE
 - DELETE CASCADE: Deleting a row in the parent table deletes all matching rows in the child table
 - UPDATE CASCADE: Updating a primary key value in the parent table updates the corresponding foreign key values in the child table
 - Useful when the child records are entirely dependent on the parent record

RELATIONAL CONSTRAINTS

- RESTRICT
 - DELETE RESTRICT: Blocks deletion of a row in the parent table if a child table has any matching foreign key references
 - UPDATE RESTRICT: Blocks updating the primary key in the parent table if a child table has any matching foreign key references
 - Enforces the integrity of the child data by ensuring it cannot be removed or updated indirectly.
- SET NULL
 - DELETE SET NULL: Deleting a row in the parent table causes the corresponding foreign key values in the child table to be set to NULL.
 - UPDATE SET NULL: Updating a primary key in the parent table causes the corresponding foreign key values in the child table to be set to NULL.
 - Used when the child record should remain but the relationship to the parent should be removed

RELATIONAL CONSTRAINTS

- SET DEFAULT

- DELETE SET DEFAULT: When a row in the parent table is deleted, the corresponding foreign key values in the child table are set to a predefined default.
- UPDATE SET DEFAULT: When a primary key in the parent table is updated, the corresponding foreign key values in the child table are set to the default value.
- Useful when a default relationship should be established if the original relationship is removed.
 - *Eg. The default sales person is the sales manager, when deleting a sales persons, all their accounts are transferred to the sales manager*

- NO ACTION

- Any attempt to delete or update a row in the parent table that is referenced by a child row will result in an error.
- Similar to RESTRICT, but the check is deferred until the end of the transaction, allowing for temporary inconsistencies within a transaction that can be fixed up during the transaction

RELATIONAL CONSTRAINTS

- RESTRICT
 - DELETE RESTRICT: Blocks deletion of a row in the parent table if a child table has any matching foreign key references
 - UPDATE RESTRICT: Blocks updating the primary key in the parent table if a child table has any matching foreign key references
 - Enforces the integrity of the child data by ensuring it cannot be removed or updated indirectly.
- CASCADE
 - DELETE CASCADE: Deleting a row in the parent table deletes all matching rows in the child table
 - UPDATE CASCADE: Updating a primary key value in the parent table updates the corresponding foreign key values in the child table
 - Useful when the child records are entirely dependent on the parent record

TRIGGERS

- Database objects that automatically execute in response to certain events
 - Help automate complex database logic, enforce constraints, maintain audit trails, and keep data consistent.
 - Fire when specific events (like INSERT, UPDATE, or DELETE) occur on a table or view.
 - Can be set to fire BEFORE or AFTER the triggering event, or even INSTEAD OF in the case of views.
 - Row-level triggers execute once for each row affected by the triggering event.
 - Statement-level triggers execute once per SQL statement, regardless of the number of rows affected.
- Triggers require a trigger function, which contains the actions or logic to be executed.
 - Written in PL/pgSQL or other supported languages.

LAB 4-2

- The lab description and documentation is in the Lab directory in the class repository



DESIGNING TABLES WELL

- Understand the Data and Its Usage Patterns
 - Use a data model to understand what the data represents and what constitutes valid data
 - Understand what will data accessed and used.
 - Consider the types of queries that will be run, how often the data will be updated, and the relationships between different entities.
- Choose Appropriate Data Types
 - Use the most suitable data types for each column to optimize storage and performance.
 - *For example, use `INTEGER` instead of `BIGINT` if the range of values allows it, or `VARCHAR(n)` with a defined length instead of `TEXT` when the length is predictable.*
 - Leverage PostgreSQL specific data types like `JSONB` for JSON data, `ARRAY` for lists, and `UUID` for unique identifiers.

DESIGNING TABLES WELL

- Normalize The Data
 - Apply normalization up to the third normal form to reduce redundancy
 - Avoid over-normalization, which can lead to excessive joins and performance issues.
 - Strike a balance between normalization and practical performance needs based on profiling performance
- Define Primary Keys
 - Ensure every table has a primary key that uniquely identifies each row.
 - Use SERIAL or BIGSERIAL types for auto-incrementing primary keys, or use UUID for distributed systems where uniqueness across databases is needed.
- Use Constraints to Enforce Data Integrity
 - Utilize constraints such as NOT NULL, UNIQUE, PRIMARY KEY, CHECK, and FOREIGN KEY to enforce rules and maintain data integrity.

DESIGNING TABLES WELL

- Use Indexes
 - Create indexes on columns that are frequently used in WHERE clauses, joins, and sorting operations to improve query performance.
 - Use appropriate index types: B-tree indexes for general-purpose indexing, GIN indexes for full-text search or JSONB data, and GiST indexes for geometric data.
 - Avoid over-indexing because indexes consume storage and can slow down write operations like INSERT, UPDATE, and DELETE.
- Partition Large Tables
 - Consider partitioning large tables to improve performance and manageability based on a partition key that is consistent with domain logic
 - Partitioning can reduce query times by scanning only relevant partitions and can improve maintenance tasks like vacuuming and backups.

DESIGNING TABLES WELL

- Optimize for Read and Write Patterns
 - Analyze how data will be accessed and optimized for the most frequent read and write operations.
 - For read-heavy workloads, consider denormalization or using materialized views to simplify and speed up complex queries.
 - For write-heavy workloads, minimize the number of indexes and avoid complex constraints that can slow down data modifications.
- Avoid NULLs Where Possible
 - Avoid excessive use of NULL values, which can complicate queries and affect indexing.
 - Use default values or redesign schemas to minimize the presence of NULLs, especially in indexed columns.

DESIGNING TABLES WELL

- Use Foreign Keys and Define Relationships Clearly
 - Use foreign keys to enforce referential integrity between tables, ensuring that relationships between data remain valid.
 - Use ON DELETE and ON UPDATE actions (CASCADE, SET NULL, RESTRICT) to specify what should happen when related data is modified.
- Data Security and Access Control
 - Plan for security by defining roles and permissions that restrict access to sensitive data.
 - Use views to provide controlled access to data and prevent direct access to base tables when necessary.
- Document the Schema
 - Maintain clear documentation for your schema, including descriptions of tables, columns, constraints, and the purpose of each.
 - Use descriptive names for tables, columns, and constraints to make the schema self-explanatory.

DESIGNING TABLES WELL

- Plan for Scalability and Growth
 - Design tables with future growth in mind, considering how the database might need to scale in terms of data volume and concurrent access.
 - Use sequences or UUIDs for keys that need to scale across distributed systems.
- Regularly Analyze and Vacuum Tables
 - Use the ANALYZE command to update statistics for the query planner, which helps optimize query execution.
 - Regularly run VACUUM to reclaim storage from deleted or updated rows and to prevent table bloat.

DESIGNING TABLES WELL

- Monitor Performance and Adjust
 - Continuously monitor database performance using PostgreSQL tools like `pg_stat_activity`, `pg_stat_user_tables`, and the `EXPLAIN` command to analyze query plans.
 - Make adjustments based on performance metrics, such as adding or removing indexes, adjusting table design, or modifying queries.

End Module



PostgreSQL