

4. Managing Tables and Views

Introduction to PostgreSQL



PostgreSQL

AGENDA

- Basic table creation
- Table constraints
- Table relationships
- Referential integrity
- Designing tables well



TABLE MANAGEMENT

- Tables and related objects are created in PostgreSQL using standard SQL commands
- Tables are created in a schema in a database
 - If no schema is specified, then the table is created in the default public schema
- Tables are owned by the user that created them
 - Access privileges are given by the owner to other users with GRANT



TABLE CREATION

- Created using CREATE TABLE
 - List tables with \dt
 - List table details with \d <tablename>
- Delete tables with DROP TABLE
 - All data in the table is lost
 - Tables can only be dropped by the owner or user with permission to drop the table

```
CREATE TABLE weather (  
    city          varchar(80),  
    temp_lo       int,          -- low temperature  
    temp_hi       int,          -- high temperature  
    prcp          real,         -- precipitation  
    date          date  
);  
  
CREATE TABLE cities (  
    name          varchar(80),  
    location       point  
);
```

```
lab4=# DROP TABLE cities, weather;  
lab4=# \dt  
        List of relations  
 Schema | Name   | Type  | Owner  
-----+-----+-----+-----  
(0 rows)
```

CREATE TABLE AS

- Create new table by copying the data and structure of the existing table.
 - WITH NO DATA only copies the structure
 - Otherwise the data is copied
- SELECT statements can be used to populate the new table with a subset of the original data

```
lab4=# SELECT * FROM employees;
```

id	name	department	salary
1	Alice	Sales	50000
2	Bob	Marketing	55000
3	Charlie	Sales	60000
4	David	IT	65000
5	Eve	HR	48000

(5 rows)

```
lab4=# CREATE TABLE ecopy AS TABLE employees;
```

```
lab4=# SELECT * FROM ecopy;
```

id	name	department	salary
1	Alice	Sales	50000
2	Bob	Marketing	55000
3	Charlie	Sales	60000
4	David	IT	65000
5	Eve	HR	48000

(5 rows)

TABLE INHERITANCE

- In the OO view, tables represent types of object
 - We can create sub-types that inherit from base types
- The sub-type
 - Has the same columns as its parent
 - Generally has more attributes or columns
- This is used for table design rather than storing normalized data.
 - Working with data across multiple child tables can be complex
 - Often partitioning is a better solution

```
CREATE TABLE cities (  
    name      text,  
    population float8,  
    elevation  int    -- (in ft)  
);  
  
CREATE TABLE capitals (  
    state      char(2)  
) INHERITS (cities);  
  
-- Now, let's populate the tables.  
INSERT INTO cities VALUES ('San Francisco', 7.24E+5, 63);  
INSERT INTO cities VALUES ('Las Vegas', 2.583E+5, 2174);  
INSERT INTO cities VALUES ('Mariposa', 1200, 1953);  
  
INSERT INTO capitals VALUES ('Sacramento', 3.694E+5, 30, 'CA');  
INSERT INTO capitals VALUES ('Madison', 1.913E+5, 845, 'WI');  
  
SELECT * FROM cities;  
SELECT * FROM capitals;
```




TABLE PARTITIONING

- Partitioning divides large tables into smaller, more manageable pieces
 - Different methods for partitioning are:
 - Range – data is grouped by ranges of values for a specific attribute, like transaction data
 - List – similar to range but on a categorical attribute, like county name for example
 - Hash – uses a hash c part partitioning methods. Here's an example of how to create a partitioned table using range partitioning:

```
CREATE TABLE cities (  
  name      text,  
  population float8,  
  elevation  int    -- (in ft)  
);  
  
CREATE TABLE capitals (  
  state      char(2)  
) INHERITS (cities);  
  
-- Now, let's populate the tables.  
INSERT INTO cities VALUES ('San Francisco', 7.24E+5, 63);  
INSERT INTO cities VALUES ('Las Vegas', 2.583E+5, 2174);  
INSERT INTO cities VALUES ('Mariposa', 1200, 1953);  
  
INSERT INTO capitals VALUES ('Sacramento', 3.694E+5, 30, 'CA');  
INSERT INTO capitals VALUES ('Madison', 1.913E+5, 845, 'WI');  
  
SELECT * FROM cities;  
SELECT * FROM capitals;
```

TABLE PARTITIONING

- Partitioning divides large tables into smaller, more manageable pieces
 - Range partitioning – data is grouped by ranges of values for a specific attribute, like transaction data
 - List – similar to range but on a categorical attribute, like county name for example
 - Hash – uses hashing into partitions based on hashing a specific attribute
 - Provides a more uniform distribution of data if the data is skewed along a partition attribute
 - We define the number of partitions
 - Useful with large data sets because it improves the performance of queries



TABLESPACES

- Tablespace define locations on the file system where the database objects are stored.
 - By default all data is stored in the default tablespace located in the data directory.
 - Tablespaces define other locations for storing data by mapping it to specific physical locations.
 - Created by defining a directory PostgreSQL has permission to read and write data files.
- Benefits
 - By distributing data across different physical disks, I/O performance can be improved.
 - New tablespaces on different disks can expand physical storage without moving the entire database.
 - Allows DBAs to specify where specific tables, indexes, or entire schemas are stored, making it easier to manage large and complex databases.
 - Provides a way to separate data physically.
 - Streamlines backup and recovery

TABLESPACES

- Possible issues
 - Can add complexity administration because of track of where different object are located.
 - For physical backups tablespaces must be backed up and restored to their respective locations or errors could result.
 - Logical backups don't capture tablespace information by default
- pg_default Tablespace
 - Where user defined database are stored by default unless specified otherwise.
- pg_global Tablespace
 - Used for storing shared system catalogs and objects that are shared across the entire cluster.
 - *Objects like pg_database, which keeps track of all databases in the cluster, and pg_authid, which stores role and authentication information. T*
 - Not accessible for user defined objects

VIEWS

- Virtual tables that represent the result of a query.
 - The query is stored in raw SQL format so it can be reused
 - Used to simplify complex queries, encapsulate business logic, and present data without duplicating the actual data in the underlying tables
- Limitations of Views:
 - Views execute their underlying query each time they are accessed
 - Complex views can have performance issues, especially if there are multiple joins or large datasets.
- Not all views are updateable
 - Updateable means that changes made to the view are propagated into the tables
 - Updateable must meet specific criteria, such as being based on a single table without groupings or aggregates.

VIEWS

- Updateable view requirements
 - The view must be defined by a simple SELECT statement that meets the following criteria:
 - *It selects from a single table.*
 - *It does not include any aggregation (GROUP BY, HAVING).*
 - *It does not include DISTINCT, set operations (UNION, INTERSECT, EXCEPT), or subqueries.*
 - *It does not include window functions.*
 - *It does not use LIMIT or OFFSET.*
 - *It does not contain JOIN clauses.*
 - Each column in the view must directly map to a column in the underlying base table without any expressions, computations, or column aliases that modify data types or values.
 - The underlying table should have a primary key or a unique constraint to uniquely identify rows.

MATERIALIZED VIEWS

- Similar to regular views
 - Unlike regular views, materialized views store the result set of the query on disk, allowing for faster access to the data.
 - Since the data is stored, materialized views need to be refreshed to reflect changes in the underlying tables.
 - Refreshing can be done manually or automatically.
 - Ideal for queries that are expensive to run and do not need real-time data accuracy, such as reporting or analytical queries.

DATA TYPES

- PostgreSQL supports the standard SQL data types
 - Also supports a number of complex data types
 - This will be reviewed in the documentation
- Default values can be specified
 - Either values of some computation
 - Like the 'now()' function



CREATING DATA TYPES

- PostgreSQL also allows user to create custom data types
- Common custom types
 - Composite Types: allow grouping multiple fields of different data types into a single logical unit.
 - Enum Types: static, predefined values for attributes with a limited set of possible values
 - *Excellent for enforcing referential integrity*
 - Domain Types: on existing data types but include constraints
 - *CREATE DOMAIN positive_int AS INT CHECK (VALUE > 0);*

```
CREATE TYPE full_name AS (  
    first_name TEXT,  
    last_name TEXT  
);  
  
CREATE TABLE employees (  
    employee_id SERIAL PRIMARY KEY,  
    name full_name,  
    position TEXT  
);  
  
INSERT INTO employees (name, position) VALUES  
(('John', 'Doe'), 'Software Developer'),  
(('Jane', 'Smith'), 'Project Manager');  
  
SELECT name.first_name FROM employees;
```


CUSTOM DATA TYPES

- Benefits
 - Maintain data integrity by enforcing consistent data structures and validation rules.
 - Encapsulating related fields into a single type allows for simpler queries
 - Promotes re-usability across multiple tables and applications, reducing redundancy.
 - User-defined types makes PostgreSQL highly extensible, for specific application needs.
- Considerations
 - Custom types can increase schema complexity
 - Some custom types, especially those requiring complex operations, might impact performance.
 - Might introduce compatibility challenges when interfacing with other systems or migrating data.

LAB 4-1

- The lab description and documentation is in the Lab directory in the class repository



TABLE CONSTRAINTS

- Table constraints are rules applied to columns and tables
 - Enforce data integrity, ensure consistency, and define the relationships between data.
 - Help prevent invalid data from being entered into the database
 - *Eg. Two customers with the same customer number*
 - Provide safeguards that maintain the logical correctness of the data.
 - *Eg. Invoice for a non-existent customer*
- Types of constraints
 - UNIQUE: ensures all values in a column or group of columns are unique across the table.
 - NOT NULL: Ensures that a column cannot have NULL values.
 - PRIMARY KEY: Combination of UNIQUE and NOT NULL
 - *Each table can have at most one primary key*
 - *Often a single column but can be combination of columns – called a composite key*

TABLE CONSTRAINTS

- FOREIGN KEY: Enforces a link between the data in two tables for referential integrity
 - *References the primary key or unique constraint of another table.*
 - *Prevents actions that would Invalidate the link like deleting the foreign table row if there are dependent rows in the table where the foreign key constraint is defined*
 - *Can specify actions on update or delete (e.g., CASCADE, SET NULL, RESTRICT)*
- CHECK: Enforces custom conditions that must be true for each row in the table.
 - *Similar to constraints when defining a domain.*
- EXCLUSION: Ensures that, for any two rows, the specified columns do not have overlapping values
 - *Commonly used with range and geometric data types.*
 - *Useful for scenarios like scheduling where no two time periods should overlap.*

TABLE CONSTRAINTS

- Constraints ensure that the data entered into a table meets specified criteria, maintaining data integrity and reliability.
- Constraints can also impact performance, especially on large tables with frequent updates.
 - Proper indexing and query optimization are necessary to mitigate performance degradation
- Constraints can be combined to enforce complex rules.
 - For instance, a column can be both NOT NULL and UNIQUE
- Some constraints, such as foreign keys, declared DEFERRED
 - This means they are checked at the end of a transaction rather than immediately

RELATIONAL CONSTRAINTS

- The primary use of a relational constraint
 - Defines what should happen to a row in a child table that has a foreign key (parent table)
 - Changes in the parent table can invalidate the relationship, not the data
 - *For example, the customer table has a reference to the sales person table. What happens when a sales person is fired?*
- The changes that can affect referential integrity are DELETE and UPDATE for the parent table
- CASCADE
 - DELETE CASCADE: Deleting a row in the parent table deletes all matching rows in the child table
 - UPDATE CASCADE: Updating a primary key value in the parent table updates the corresponding foreign key values in the child table
 - Useful when the child records are entirely dependent on the parent record

RELATIONAL CONSTRAINTS

- RESTRICT
 - DELETE RESTRICT: Blocks deletion of a row in the parent table if a child table has any matching foreign key references
 - UPDATE RESTRICT: Blocks updating the primary key in the parent table if a child table has any matching foreign key references
 - Enforces the integrity of the child data by ensuring it cannot be removed or updated indirectly.
- SET NULL
 - DELETE SET NULL: Deleting a row in the parent table causes the corresponding foreign key values in the child table to be set to NULL.
 - UPDATE SET NULL: Updating a primary key in the parent table causes the corresponding foreign key values in the child table to be set to NULL.
 - Used when the child record should remain but the relationship to the parent should be removed

RELATIONAL CONSTRAINTS

- SET DEFAULT

- DELETE SET DEFAULT: When a row in the parent table is deleted, the corresponding foreign key values in the child table are set to a predefined default.
- UPDATE SET DEFAULT: When a primary key in the parent table is updated, the corresponding foreign key values in the child table are set to the default value.
- Useful when a default relationship should be established if the original relationship is removed.
 - *Eg. The default sales person is the sales manager, when deleting a sales persons, all their accounts are transferred to the sales manager*

- NO ACTION

- Any attempt to delete or update a row in the parent table that is referenced by a child row will result in an error.
- Similar to RESTRICT, but the check is deferred until the end of the transaction, allowing for temporary inconsistencies within a transaction that can be fixed up during the transaction

RELATIONAL CONSTRAINTS

- RESTRICT
 - DELETE RESTRICT: Blocks deletion of a row in the parent table if a child table has any matching foreign key references
 - UPDATE RESTRICT: Blocks updating the primary key in the parent table if a child table has any matching foreign key references
 - Enforces the integrity of the child data by ensuring it cannot be removed or updated indirectly.
- CASCADE
 - DELETE CASCADE: Deleting a row in the parent table deletes all matching rows in the child table
 - UPDATE CASCADE: Updating a primary key value in the parent table updates the corresponding foreign key values in the child table
 - Useful when the child records are entirely dependent on the parent record

TRIGGERS

- Database objects that automatically execute in response to certain events
 - Help automate complex database logic, enforce constraints, maintain audit trails, and keep data consistent.
 - Fire when specific events (like INSERT, UPDATE, or DELETE) occur on a table or view.
 - Can be set to fire BEFORE or AFTER the triggering event, or even INSTEAD OF in the case of views.
 - Row-level triggers execute once for each row affected by the triggering event.
 - Statement-level triggers execute once per SQL statement, regardless of the number of rows affected.
- Triggers require a trigger function, which contains the actions or logic to be executed.
 - Written in PL/pgSQL or other supported languages.

LAB 4-2

- The lab description and documentation is in the Lab directory in the class repository



DESIGNING TABLES WELL

- Understand the Data and Its Usage Patterns
 - Use a data model to understand what the data represents and what constitutes valid data
 - Understand what will data accessed and used.
 - Consider the types of queries that will be run, how often the data will be updated, and the relationships between different entities.
- Choose Appropriate Data Types
 - Use the most suitable data types for each column to optimize storage and performance.
 - *For example, use `INTEGER` instead of `BIGINT` if the range of values allows it, or `VARCHAR(n)` with a defined length instead of `TEXT` when the length is predictable.*
 - Leverage PostgreSQL specific data types like `JSONB` for JSON data, `ARRAY` for lists, and `UUID` for unique identifiers.

DESIGNING TABLES WELL

- Normalize The Data
 - Apply normalization up to the third normal form to reduce redundancy
 - Avoid over-normalization, which can lead to excessive joins and performance issues.
 - Strike a balance between normalization and practical performance needs based on profiling performance
- Define Primary Keys
 - Ensure every table has a primary key that uniquely identifies each row.
 - Use SERIAL or BIGSERIAL types for auto-incrementing primary keys, or use UUID for distributed systems where uniqueness across databases is needed.
- Use Constraints to Enforce Data Integrity
 - Utilize constraints such as NOT NULL, UNIQUE, PRIMARY KEY, CHECK, and FOREIGN KEY to enforce rules and maintain data integrity.

DESIGNING TABLES WELL

- Use Indexes
 - Create indexes on columns that are frequently used in WHERE clauses, joins, and sorting operations to improve query performance.
 - Use appropriate index types: B-tree indexes for general-purpose indexing, GIN indexes for full-text search or JSONB data, and GiST indexes for geometric data.
 - Avoid over-indexing because indexes consume storage and can slow down write operations like INSERT, UPDATE, and DELETE.
- Partition Large Tables
 - Consider partitioning large tables to improve performance and manageability based on a partition key that is consistent with domain logic
 - Partitioning can reduce query times by scanning only relevant partitions and can improve maintenance tasks like vacuuming and backups.

DESIGNING TABLES WELL

- Optimize for Read and Write Patterns
 - Analyze how data will be accessed and optimized for the most frequent read and write operations.
 - For read-heavy workloads, consider denormalization or using materialized views to simplify and speed up complex queries.
 - For write-heavy workloads, minimize the number of indexes and avoid complex constraints that can slow down data modifications.
- Avoid NULLs Where Possible
 - Avoid excessive use of NULL values, which can complicate queries and affect indexing.
 - Use default values or redesign schemas to minimize the presence of NULLs, especially in indexed columns.

DESIGNING TABLES WELL

- Use Foreign Keys and Define Relationships Clearly
 - Use foreign keys to enforce referential integrity between tables, ensuring that relationships between data remain valid.
 - Use ON DELETE and ON UPDATE actions (CASCADE, SET NULL, RESTRICT) to specify what should happen when related data is modified.
- Data Security and Access Control
 - Plan for security by defining roles and permissions that restrict access to sensitive data.
 - Use views to provide controlled access to data and prevent direct access to base tables when necessary.
- Document the Schema
 - Maintain clear documentation for your schema, including descriptions of tables, columns, constraints, and the purpose of each.
 - Use descriptive names for tables, columns, and constraints to make the schema self-explanatory.

DESIGNING TABLES WELL

- Plan for Scalability and Growth
 - Design tables with future growth in mind, considering how the database might need to scale in terms of data volume and concurrent access.
 - Use sequences or UUIDs for keys that need to scale across distributed systems.
- Regularly Analyze and Vacuum Tables
 - Use the ANALYZE command to update statistics for the query planner, which helps optimize query execution.
 - Regularly run VACUUM to reclaim storage from deleted or updated rows and to prevent table bloat.

DESIGNING TABLES WELL

- Monitor Performance and Adjust
 - Continuously monitor database performance using PostgreSQL tools like `pg_stat_activity`, `pg_stat_user_tables`, and the `EXPLAIN` command to analyze query plans.
 - Make adjustments based on performance metrics, such as adding or removing indexes, adjusting table design, or modifying queries.

End Module



PostgreSQL