

5. Performance Optimization

Introduction to PostgreSQL



PostgreSQL

AGENDA

- SQL raw storage format
- The query planner
- Queries and code optimization
- Indexing
- Schema design



RAW SQL STORAGE

- For certain database objects like views, rules, and functions
 - PostgreSQL stores the original SQL text used to define these objects as raw SQL strings in specific system catalog tables.
- This raw SQL text is preserved so that PostgreSQL can reproduce the object definitions
- For many objects, the raw SQL is not stored, tables for example because
 - Storing tables as structured metadata allows faster access and managing of table information without parsing raw SQL every time.
 - Changes to table structures, such as adding columns or modifying constraints, can be efficiently managed at the metadata level without the need to re-parse or reconstruct SQL.
 - Managing table structures through metadata ensures that all dependencies and internal data structures are consistently maintained.

RAW SQL STORAGE

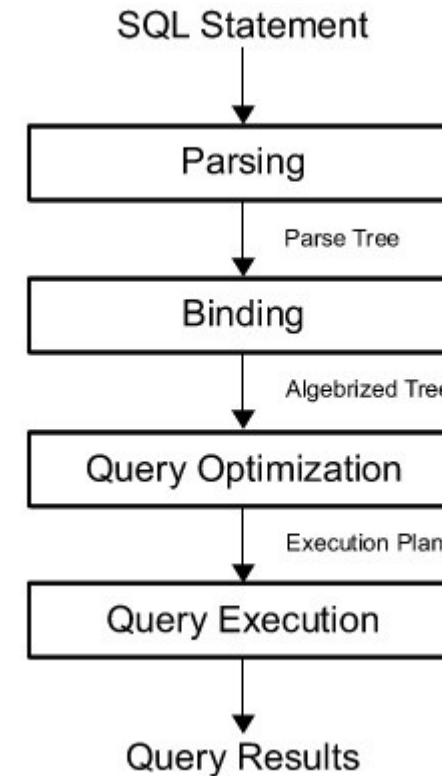
- Why use raw SQL storage?
 - Primarily for object not stored like tables in an internal schema format
 - Tools like pg_dump rely on raw SQL definitions to accurately some objects when exporting that when a database is restored or replicated, the objects are recreated exactly as they were originally defined.
 - Objects like views, rules, and functions often contain complex logic is straightforward to reference, modify, recompile or duplicate from the raw SQL.

RAW SQL STORAGE

- What is not in raw SQL – some examples
- Table definitions
- Indexes
 - Indexes are highly optimized data structures that rely on detailed internal representation to support efficient lookups.
- Constraints and Triggers
 - Storing constraints as metadata, allows enforcing data integrity directly at the storage level without needing to refer back to SQL definitions.
- DML Commands (INSERT, UPDATE, DELETE, SELECT)
 - DML commands are intended to manipulate data and do not define persistent database objects, so there's no need to store their SQL.


QUERY OPTIMIZATION

- Query optimizers analyze SQL queries and find the best execution plan
- Best means minimizing the resources used, such as time, memory, and CPU
- And minimizing the time return a result
- This is a standard component of any SQL or similar system that queries data
- For example
 - A join on 100,000 records followed by a select that only keep 1,000 records
 - Reversing the order of operations means much less data is used in the join




OPTIMIZATION STEPS

- Parsing the Query:
 - SQL is parsed into a tree to understand its structure and what data is being requested.
 - This step includes checking for syntax errors and creating a preliminary plan called a parse tree.
- Generating Possible Execution Plans
 - The optimizer generates several possible logical execution plans.
 - The plans try different methods to access and join tables, such as using indexes, performing full table scans, or joining tables in different orders.
 - The query may be rewritten to make it more efficient



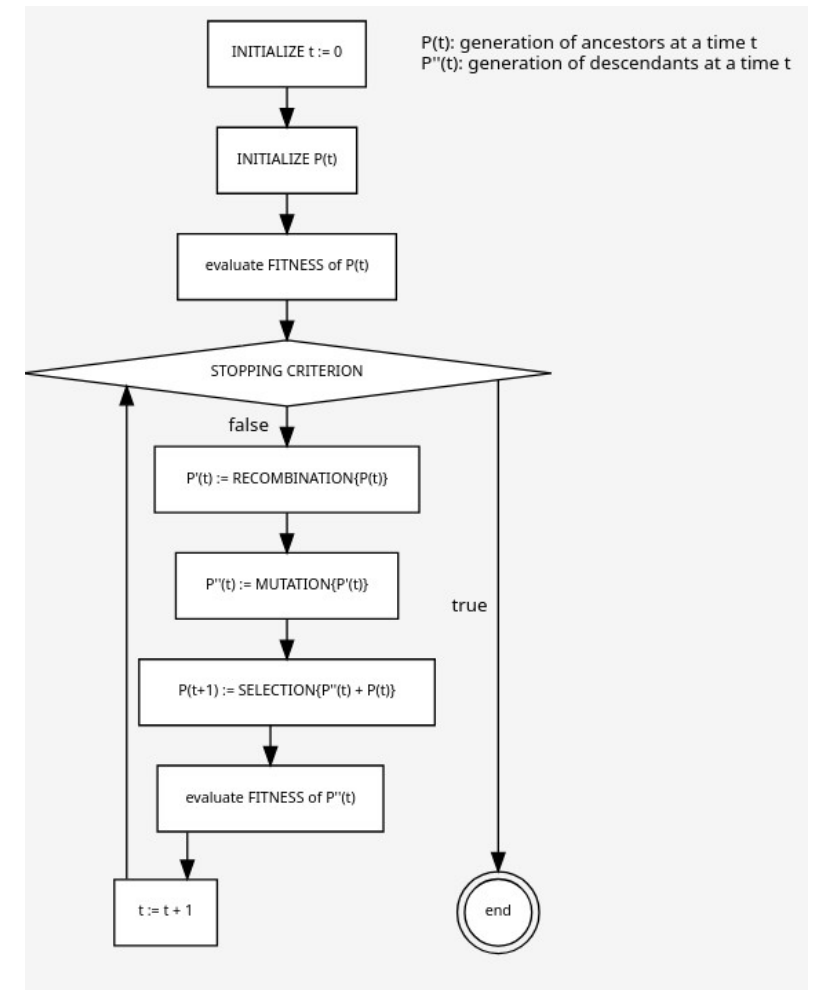
THE JOIN PROBLEM

- The problem of efficiently executing SQL queries that involve joins between multiple tables
 - As the number of tables in a join increases, the number of possible join orders grows at a factorial rate ($x!$ In the diagram)
 - With three tables, there are only a few join orders to consider, but with ten tables, there are over three million possible join combinations.
 - Many potential join orders need to be considered to find the most efficient one
 - This can become computationally expensive and time-consuming



GENETIC QUERY OPTIMIZATION (GEQO)

- The GEQO planning process uses the standard planner code to generate plans for scans of individual relations.
 - Then join plans are developed using the genetic approach.
 - This process is inherently nondeterministic, because of the randomized choices made during both the initial population selection and subsequent “mutation” of the best candidates.



SQL QUERY PLANNING

- The application program transmits a query to the server and waits to receive the results sent back by the server.
- The parser stage checks the query transmitted by the application program for correct syntax and creates a query tree.
- The rewrite system takes the query tree created by the parser stage and looks for any rules (stored in the system catalogs) to apply to the query tree.
 - It performs the transformations given in the rule bodies.
 - One application of the rewrite system is in the realization of views. Whenever a query against a view (i.e., a virtual table) is made, the rewrite system rewrites the user's query to a query that accesses the base tables given in the view definition instead.
- The planner/optimizer takes the (rewritten) query tree and creates a query plan that will be the input to the executor.
 - It does so by first creating all possible paths leading to the same result.
 - For example if there is an index on a relation to be scanned, there are two paths for the scan.
 - One possibility is a simple sequential scan and the other possibility is to use the index.
 - Next the cost for the execution of each path is estimated and the cheapest path is chosen. The cheapest path is expanded into a complete plan that the executor can use.

SQL QUERY PLANNING

- The executor recursively steps through the plan tree and retrieves rows in the way represented by the plan.
 - The executor makes use of the storage system while scanning relations, performs sorts and joins, evaluates qualifications and finally hands back the rows derived.

DEEP DIVE

- This is an optional deep dive into the PostgreSQL rule system



EXPLAIN

- The EXPLAIN command shows the steps that the database will take to retrieve or modify the data as requested by the query.
 - Using EXPLAIN can identify potential performance issues and optimize queries accordingly.
- What EXPLAIN does?
 - Shows the execution plan chosen by the query optimizer, including the sequence of operations (e.g., scans, joins, sorts) and their order.
 - Shows estimated costs for each operation, including startup and total costs, which reflect resource usage (like CPU and I/O)
 - Provides estimates of the number of rows processed by each step.
 - When used with ANALYZE, it provides the actual execution times and row counts, which are useful for comparing the optimizer's estimates with real performance.
- The lab works with an actual example.

LAB 5-1

- The lab description and documentation is in the Lab directory in the class repository



OPTIMIZATION TIPS

- Use Proper Indexing:
 - Create Indexes on frequently queried columns: Index columns used in WHERE clauses, joins, and ORDER BY clauses.
 - Use Composite Indexes: When multiple columns are frequently queried together, consider composite indexes.
 - Use Partial Indexes: Index only the relevant subset of rows to reduce index size and improve performance.
 - Avoid Over-Indexing: Too many indexes can slow down INSERT, UPDATE, and DELETE operations.
- Optimize Joins:
 - Choose the Right Join Type: Use inner joins for matching rows and outer joins only when necessary.
 - Use Indexed Columns in Joins: Ensure join columns are indexed to speed up join operations.
 - Reduce Join Complexity: Simplify joins by breaking down complex queries or reducing the number of tables involved.

OPTIMIZATION TIPS

- Leverage Query Planning and Execution Tools:
 - Use EXPLAIN and EXPLAIN ANALYZE: Review execution plans to understand how queries are executed and identify performance bottlenecks.
 - Check for Sequential Scans: Ensure that large table scans are necessary and consider indexing to avoid them.
- Write Efficient SQL:
 - Avoid SELECT *: Specify only the columns you need to reduce I/O and processing time.
 - Use Subqueries and CTEs Judiciously: While useful, complex subqueries and Common Table Expressions (CTEs) can sometimes be replaced with simpler joins or views.
 - Avoid Unnecessary Calculations: Avoid repetitive calculations within queries by computing values once or using derived columns.

OPTIMIZATION TIPS

- Use Appropriate Data Types:
 - Choose the Right Data Type: Use appropriate data types for columns to optimize storage and performance.
 - Avoid Over-Sized Data Types: Use INTEGER instead of BIGINT when appropriate, and avoid using large text types if not necessary.
- Optimize Filtering and Sorting:
 - Push Down Filters: Apply filters as early as possible to reduce the data set size in intermediate steps.
 - Use Indexes for Sorting: Ensure that sorting operations (ORDER BY) can utilize indexes when possible.

OPTIMIZATION TIPS

- Improve Transaction and Locking Strategies:
 - Minimize Lock Contention: Use shorter transactions and minimize the scope of locks to reduce contention.
 - Use Appropriate Isolation Levels: Choose the correct transaction isolation level (READ COMMITTED, REPEATABLE READ, etc.) based on application requirements.
- Optimize Updates and Inserts:
 - Batch Inserts and Updates: Use bulk operations instead of individual row operations to reduce overhead.
 - Use COPY for Bulk Loads: The COPY command is much faster than INSERT for loading large volumes of data.

OPTIMIZATION TIPS

- Tune Memory and Configuration Settings:
 - Adjust `work_mem`: Set appropriate `work_mem` for complex queries to allow enough memory for sorting and joining.
 - Set `shared_buffers`: Allocate sufficient memory to `shared_buffers` to store frequently accessed data in memory.
 - Configure `maintenance_work_mem`: Set a higher value for maintenance operations like vacuuming and indexing.
- Regular Maintenance:
 - Vacuum and Analyze Regularly: Run `VACUUM` and `ANALYZE` to update statistics and reclaim space.
 - Reindex as Needed: Periodically reindex tables to maintain index efficiency, especially after large data changes.

OPTIMIZATION TIPS


- Partition Large Tables:
 - Use Table Partitioning: Divide large tables into smaller, manageable partitions to improve query performance and manageability.
- Leverage Materialized Views:
 - Use Materialized Views: Precompute and store results of complex queries that do not change frequently to speed up access.
- Optimize Network and I/O:
 - Reduce Network Overhead: Minimize data transferred over the network by fetching only necessary rows and columns.
 - Use Connection Pooling: Use connection pooling to reduce the overhead of establishing database connections.

OPTIMIZATION TIPS

- Avoid Using Loops in PL/pgSQL:
 - Use Set-Based Operations: Prefer set-based SQL operations over looping row by row in PL/pgSQL functions.
- Monitor and Profile Queries:
 - Use PostgreSQL Monitoring Tools: Utilize tools like pg_stat_statements, pgBadger, or other performance monitoring solutions to identify slow queries and resource bottlenecks.


INDEXES

- PostgreSQL has several types of indexes, each suited to different use cases and query patterns
- B-Tree index
 - Are the default and most commonly used index type in PostgreSQL.
 - Work well for equality (=), range (<, >, BETWEEN), and pattern matching (LIKE) queries with non-wildcard patterns
 - Optimal for primary key, unique constraints, and general-purpose indexing on columns used in WHERE, ORDER BY, and join conditions.



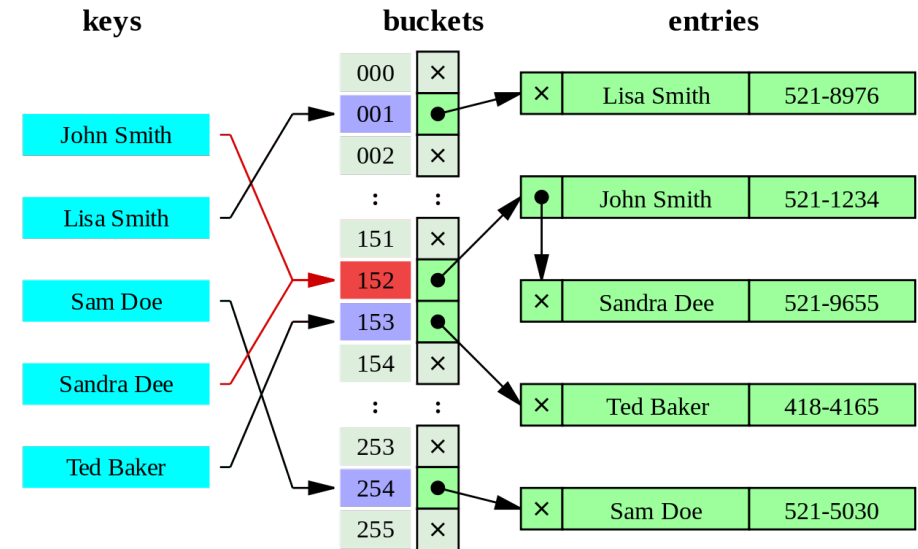
INDEXES

- Performance Tips:
 - Use B-tree indexes on frequently queried columns with a high level of distinct values.
 - Avoid using B-tree indexes on columns with many repeated values which will impact performance.



INDEXES

- Hash Index
 - Hash indexes are used specifically for equality comparisons (=).
 - Generally faster than B-tree indexes for equality lookups but do not support range queries.
 - Ideal for indexing columns that are frequently queried with equality conditions.
- Performance Tips:
 - Use hash indexes when the primary use case involves equality searches.
 - Hash indexes require additional maintenance operations (e.g., REINDEX) after crashes, as they are not WAL-logged.



INDEXES

- Generalized Inverted Index
 - Useful for containment queries (@>, <@) and for searching elements within arrays or documents.
 - Optimal for full-text search (tsvector), JSONB data types, and columns storing arrays or other composite data types.
 - Use GIN indexes when querying JSONB, array elements, or for full-text search scenarios.
 - Be mindful of the higher maintenance cost compared to B-tree indexes.

INDEXES

- Generalized Search Tree
 - GiST indexes support a wide range of queries, including geometric data types (e.g., points, polygons), full-text search, and custom data types with complex queries.
 - Use for nearest-neighbor searches, spatial data, and range queries on complex data types.
 - Use GiST indexes for spatial and geometric queries.
 - Custom operator classes allow GiST indexes to be used for various complex data types, but they may have higher space usage.

INDEXES

- Block Range Index
 - BRIN indexes are compact and efficient for indexing large, sequential data sets.
 - They work by summarizing data within physical blocks, making them suitable for columns with ordered or clustered data.
 - Best for very large tables where columns exhibit a natural ordering (e.g., timestamps, sequential IDs).
 - Use BRIN indexes for very large datasets with naturally ordered data.
 - They are space-efficient and have low maintenance costs but are less precise than B-tree indexes.

PARTIAL INDEXES

- Partial indexes are indexes that cover only a subset of rows based on a specified condition. They can be any of the previous index types but with a WHERE clause to limit indexed rows.
 - Ideal when only a subset of the data is frequently queried.
 - `CREATE INDEX idx_active_employees ON employees (name) WHERE active = true;`
 - Use partial indexes to reduce index size and maintenance by focusing on the most relevant data.
- Ensure the index condition matches frequently queried conditions to maximize performance benefits.

PERFORMANCE OPTIMIZATION

- Keep Statistics Updated:
 - Regularly run ANALYZE to update table and index statistics so the query planner can make informed decisions.
- Monitor Index Usage:
 - Use tools like pg_stat_user_indexes to monitor index usage and identify unused or redundant indexes.
- Balance Index Overhead:
 - Avoid excessive indexing, which can slow down write operations (inserts, updates, deletes).
- Use EXPLAIN to Analyze Queries: Regularly use EXPLAIN or EXPLAIN ANALYZE to check if your queries are using indexes effectively

SCHEMA OPTIMIZATION

- Schema optimization is designing and refining the database schema to improve the performance, scalability, and maintainability of the database.
- Normalization and Denormalization:
 - Normalization involves organizing tables and columns to reduce data redundancy and improve data integrity by dividing large tables into smaller, related tables.
 - Denormalization pre-joins often joined tables to reduce the number of joins required in queries.
- Proper Indexing:
 - Identifying which columns need indexing (e.g., primary keys, foreign keys, frequently queried columns) and choosing the right type of index (e.g., B-tree, hash, GIN, GiST).
 - Over-indexing can degrade performance on write operations (INSERT, UPDATE, DELETE) which might make indexes stale and need to be recreated.
- Data Types and Constraints:
 - Use data types for columns can save storage space and improve query performance.
 - For example, using INTEGER instead of BIGINT where appropriate, or VARCHAR(50) instead of TEXT when lengths are predictable.
 - Constraints (e.g., NOT NULL, UNIQUE, CHECK) ensure data integrity and can sometimes optimize query execution plans.

SCHEMA OPTIMIZATION

- Partitioning:
 - Dividing large tables into smaller, more manageable pieces along some value like date range improves query performance by limiting the amount of data that needs to be scanned.
- Foreign Keys and Relationships
 - Properly defining foreign keys and relationships between tables ensures data consistency and helps the optimizer understand the relationships between tables, which can influence query planning.
 - In some high-performance scenarios, foreign keys might be omitted to avoid the overhead of maintaining referential integrity, but this requires careful manual management.
- Schema Design for Query Patterns
 - Understanding common query patterns can guide schema design.
 - For example, if a table is frequently joined on a specific column, ensuring that column is indexed and optimizing the data model for those joins is crucial.
 - Designing tables with consideration for read-heavy vs. write-heavy workloads can also lead to optimizations specific to the expected load.

SCHEMA OPTIMIZATION

- Avoid Over-Engineering:
 - Use simple, clear schema rather than an overly complex one.
 - Over-engineering (e.g., too many tables, unnecessary normalization) can lead to complicated joins and maintenance challenges.
- Using Materialized Views:
 - For frequently accessed, complex queries, using materialized views can improve performance.
 - These are precomputed views stored on disk that can be refreshed periodically.

LAB 5

- The lab materials and instructions are in the repository



End Module



PostgreSQL