

8. High Availability and Replication.

Introduction to PostgreSQL



PostgreSQL

HIGH AVAILABILITY

- High Availability
 - Minimizing downtime
 - Ensuring that the database service remains accessible even during failures.
 - Often related to how a cluster hands loads at peak times
- Failover and Replication
 - Replication ensures data is copied to multiple instances
 - Failover automatically switches to a standby instance if the primary fails.
 - Automated failover tools like Patroni, repmgr, PgPool-II, and PostgreSQL Automatic Failover (PAF) can automate failover and manage the high availability setup
- Load Balancing:
 - Balancing read queries across replicas helps distribute the load, enhancing performance and availability.
 - The highest demands on the system tend to be read requests.

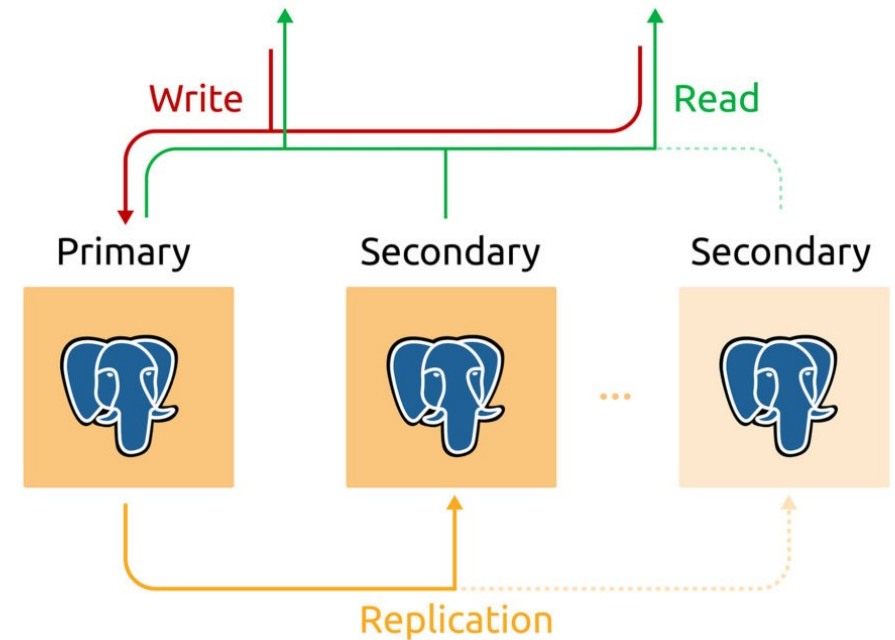
REPLICATION

- Replication
 - Copying data from one database (primary) to others (replicas).
- Streaming Replication:
 - Uses WAL (Write-Ahead Logging) to stream changes from the primary to replicas in near real-time.
 - Replicas can be set up as hot standbys (readable) or warm standbys (not readable until promoted).
- Log Shipping:
 - The primary server periodically ships completed WAL files to the standby servers, which apply them to their copy of the database



REPLICATION

- Logical Replication:
 - Replicates data changes based on logical changes (like inserts, updates, deletes)
 - Allows selective data replication, supports bi-directional replication, and is useful for replicating only specific tables or data subsets.
- Synchronous Replication:
 - Ensures that transactions are committed to both the primary and at least one standby before being acknowledged to the client.
- Asynchronous Replication:
 - Primary does not wait for the replica's acknowledgment



PLANNING

- Create the primary and standby servers so that they are as similar as possible, at least from the perspective of the database server.
 - Path names associated with tablespaces will be passed across unmodified, so both primary and standby servers must have the same mount paths for tablespaces.
 - Hardware need not be exactly the same
 - Experience shows that maintaining two identical systems is easier than maintaining two dissimilar ones over the lifetime of the application and system.
 - The hardware architecture must be the same replicating a 32-bit system to a 64-bit system will not work.
- The servers should be running the same major version
 - WAL log shipping between servers running different major releases is not possible.
 - Different minor release versions probably will work, but should be avoided

STANDBY SERVER OPERATION

- A server is in standby mode if a standby.signal file exists in the data directory when the server is started.
- In standby mode:
 - The server continuously applies WALs received from the primary server
 - Standby mode is exited and the server switches to normal operation when pg_ctl promote is run, or pg_promote() is called.
 - Before failover, any WAL immediately available in the archive or in pg_wal will be restored, but no attempt is made to connect to the primary.

SETTING UP REPLICATION

- Configuration on the Primary Server

- Edit `postgresql.conf`:
 - Set `wal_level` to `replica` or `logical` (for logical replication).
 - Set `max_wal_senders` to the number of standby servers
 - Set `archive_mode` to `on` and configure `archive_command` if using log shipping.
 - Configure `listen_addresses` to allow connections from standby servers.
- Edit `pg_hba.conf`:
 - Add entries to allow replication connections from the standby servers.

- Configuration on the Standby Server

- Base Backup:
 - Use `pg_basebackup` or other methods to create a copy of the primary server's data directory.
 - Edit `postgresql.conf` with settings like `primary_conninfo`, pointing to the primary server.
 - `primary_conninfo` specifies the connection details for the primary server, including host, port, user, and password

REPLICATION PROCESS

- Once configured, the standby server connects to the primary server and starts streaming WAL records.
- The standby server replays these records to keep its copy of the database updated with the primary server.
- In streaming replication, this process is continuous, with the standby receiving changes as they occur on the primary.

FAILOVER

- In case of a primary server failure
 - A standby can be promoted to act as the new primary.
 - This can be done manually or automatically using tools like `pg_auto_failover` or Patroni.
- Switchover
 - A planned switch where the roles of primary and standby are reversed, usually for maintenance or load balancing.

MONITORING REPLICATION

- Several tools and views for monitoring replication,
 - `pg_stat_replication` on the primary server
 - `pg_stat_wal_receiver` on standby servers.
 - These views provide information on replication status, lag, and connection details.

MONITORING REPLICATION

- Several tools and views for monitoring replication,
 - `pg_stat_replication` on the primary server
 - `pg_stat_wal_receiver` on standby servers.
 - These views provide information on replication status, lag, and connection details.
- Handling Replication Lag
 - Replication lag can occur due to network issues, slow disk I/O, or high write activity on the primary.
 - Monitor and tune parameters like `max_standby_streaming_delay` to manage how standbys handle replication delays.
- CAP Theorem
 - Find the balance of consistency and availability needed for specific use cases

End Module



PostgreSQL